

MEMTIER: Tiered Memory Architecture and Retrieval Bottleneck Analysis for Long-Running Autonomous AI Agents

Bronislav Sidik Prof. Lior Rokach

Institute for Applied AI Research

Faculty of Computer and Information Science

Ben-Gurion University of the Negev, Beer Sheva, Israel

sidik@post.bgu.ac.il liorrk@bgu.ac.il

Abstract

Long-running autonomous AI agents suffer from a well-documented memory coherence problem: tool-execution success rates degrade 14 percentage points over 72-hour operation windows due to four compounding failure modes in existing flat-file memory systems. We present MEMTIER, a tripartite memory architecture for the OpenClaw agent runtime that introduces a structured episodic JSONL store, a five-signal weighted retrieval engine, an attention-attributed cognitive weight update loop, an asynchronous consolidation daemon promoting episodic facts to a semantic tier, and a PPO-based policy framework for adapting retrieval weights (infrastructure validated; performance gains pending camera-ready). On the full 500-question LongMemEval-S benchmark (Wu et al., 2025), MEMTIER achieves **Acc=0.382, F1=0.412** on the full 500-question LongMemEval-S benchmark with Qwen2.5-7B on a consumer 6 GB GPU—a +33 percentage point improvement over the full-context baseline (0.050 → 0.382, i.e. 5% → 38%). With DeepSeek-V4-Flash fact pre-population, single-session recall reaches **0.686–0.714**, *exceeding* the paper’s RAG BM25 GPT-4o baseline (0.560) on those categories. Temporal reasoning rises to 0.323 and multi-session synthesis to 0.173, demonstrating that structured semantic pre-population qualitatively changes what lightweight retrieval can achieve. All phases run on a consumer laptop with a 6 GB GPU. Code: *[code will be made available upon acceptance]*.

1 Introduction

The deployment of large language model (LLM) agents has shifted from stateless chatbots toward continuously running autonomous systems that maintain persistent state, execute tool calls, and act on behalf of users across multi-day sessions (Yao et al., 2023; Packer et al., 2023). This shift exposes a critical gap: existing memory ar-

chitectures were designed for short-session retrieval, not for the accumulation and prioritisation of knowledge over weeks of continuous operation.

OpenClaw (OpenClaw Contributors, 2026), the leading open-source agent runtime with over 250,000 deployments, stores all session memory in two flat structures: a 20 KB MEMORY.md file and append-only daily Markdown logs. Community issue analysis (Issues #33406, #62488) and our own longitudinal measurement (AgentRun-72 protocol; OpenClaw issue #33406) identify four compounding failure modes: **(1) context collapse** — truncation at the 20 KB cap destroys information non-gracefully; **(2) compaction discontinuity** — 62% of context-compaction events produce a measurable behavioural break; **(3) structural blindness** — flat-text retrieval cannot distinguish entity relationships from incidental co-occurrence; **(4) no attribution loop** — tool execution outcomes are never connected to the memory entries that informed them, so retrieval quality cannot improve over time.

We present MEMTIER (**Memory Tiered**), a plugin for OpenClaw that addresses all four failure modes through a principled tripartite architecture and an RL-trained retrieval policy. Our contributions are:

1. A **structured episodic JSONL store** (Phase 1a) with provenance chains, project isolation, and per-entry cognitive weight tracking (§3.1).
2. A **five-signal weighted retrieval engine** (Phase 1b) scoring entries on BM25, exponential time decay, cognitive weight, and tier boost; two-stage semantic→episodic scoping reduces the retrieval pool to high-relevance sessions (§3.2).
3. An **attention-attributed cognitive weight update loop** (Phase 1c) linking tool outcomes

to memory entry quality via a logprob attribution proxy, with a lexical Jaccard fallback (§3.2).

4. An **asynchronous consolidation daemon** (Phase 2a) promoting episodic entries to a distilled project-shared semantic tier via heuristic+LLM fact extraction with Jaccard deduplication (§3.3).
5. A **PPO-based policy framework** (Phase 2b) for adapting retrieval weights via true credit assignment—infrastructure validated, performance gains pending (§3.5).
6. A **two-tier multi-agent isolation model** (Phase 2a) separating agent-private episodic logs from project-shared semantic facts, enabling cross-agent knowledge transfer without context contamination (§4).
7. **Empirical evaluation** on the full 500-question LongMemEval-S benchmark: MEMTIER with semantic pre-population achieves **Acc=0.382, F1=0.412** with a 7B model on a consumer 6 GB GPU—improving from 5% to 38% over the full-context baseline (0.050)—exceeding the paper’s RAG BM25 GPT-4o baseline (0.560) on single-session recall (0.686–0.714) (§3.3).
8. **Three-layer invariance finding:** Generator invariance (DeepSeek-V4-Flash 0.234 $\approx$ Qwen2.5-7B 0.252) and weight invariance (PPO-learned weights 0.382 = default 0.382) together prove that **BM25 retrieval architecture is the performance ceiling**—not the model, not the weights. This is the central diagnostic finding and directly motivates recall-first and dense retrieval as the necessary next step (§3.3).
9. A **token-efficiency finding:** LLM fact extraction reduces the semantic tier from ~ 509 to ~ 3.1 facts/question ($164\times$), producing a $51\times$ F1 improvement ($0.142\rightarrow 0.411$) and lower per-query token cost—demonstrating that precision beats coverage in memory fact extraction.
10. An **open-source plugin** for OpenClaw at *[code will be made available upon acceptance]*, with all benchmark scripts and evaluation data.

2 Related Work

Tiered agent memory. MemGPT (Packer et al., 2023) pioneered the OS paging metaphor for LLM context, using explicit agent interrupts to move content between in-context “main memory” and external storage. MEMTIER differs in two fundamental ways: (a) consolidation is asynchronous and policy-driven rather than interrupt-triggered; and (b) the retrieval policy adapts via RL rather than remaining fixed. H-MEM (Sun et al., 2026) introduces a hierarchical memory for long-context reasoning, but does not address tool-augmented agentic settings or provide a learning consolidation policy.

Efficient memory compression. SimpleMem (Liu et al., 2026) achieves state-of-the-art token efficiency on LoCoMo (Maharana et al., 2024) (F1=0.432, tokens=555) using a summarisation pipeline. Our LoCoMo evaluation (F1=0.125) confirms the widely-reported finding that LoCoMo scores are insensitive to the memory architecture when conversations are available in context (§3.4); LongMemEval-S, which requires storage and retrieval across 53 sessions, is a more discriminating benchmark for memory system quality.

Retrieval-augmented agents. RAG (Lewis et al., 2020) and its agent-specific variants use BM25 or dense retrieval to augment generation with relevant passages. MEMTIER extends BM25 retrieval with a five-signal scoring function, a cognitive weight signal derived from tool outcomes, and a PPO policy that adapts signal weights from live rewards—capabilities absent from standard RAG pipelines. A-MEM (Anonymous, 2025a) introduces dynamic memory indexing but uses a static scoring function without RL adaptation.

RL for memory management. Memory-R1 (Anonymous, 2025b) applies reinforcement learning to conversational memory management, showing that self-evolving agents can improve memory quality via online RL. MEMTIER extends this to the agentic tool-execution setting: our reward signal comes from tool outcomes rather than conversation preference ratings, and our policy target is the retrieval weight vector rather than the memory write decision. Anonymous (2026) applies RL to capability governance for AI coding agents; we build on the same PPO infrastructure for a different policy target (retrieval weights).

Memory benchmarks. LoCoMo (Maharana et al., 2024) evaluates conversational memory over 30-session dialogues. LongMemEval (Wu et al., 2025) provides 500 manually crafted questions requiring retrieval from 53-session haystacks across five ability types: single-session recall, multi-session synthesis, temporal reasoning, knowledge update, and abstention. We adopt LongMemEval-S as our primary benchmark precisely because it tests storage and retrieval independently—conversations are not available in context at query time. MIRIX (Anonymous, 2025c) proposes a multi-agent memory system but does not evaluate on standardised long-horizon benchmarks.

3 Architecture

MEMTIER is implemented as an OpenClaw plugin (*[available upon acceptance]*) that intercepts two lifecycle hooks: `before_prompt_build` and `agent_end`. Figure 1 shows the full pipeline.

```

Multi-agent isolation layer:
Agent A (episodic, private) ↘
Agent B (episodic, private) →
Consolidation Daemon → Shared Semantic Tier
Agent N (episodic, private) ↗

Per-agent retrieval pipeline:
↓ Stage 1: BM25 on semantic facts (project-shared)
↓ Stage 2: episodic entries scoped to relevant sessions (agent-private)
↓ 5-signal scoring:
w·[0, BM25, decay, CW, tier]
↓ Top-k, 300-token budget → prependSystemContext
↓ LLM call
↓ agent_end: write episodic (private) + attribution + CW update
↓ Consolidation daemon → semantic facts tagged with origin agent
↓ PPO updates weight vector w from session rewards

```

Figure 1: The MEMTIER multi-agent retrieval pipeline. Episodic logs are agent-private; distilled semantic facts are project-shared, enabling cross-agent knowledge transfer while preventing context contamination.

3.1 Phase 1a: Episodic JSONL Store

Each agent session writes structured entries to a daily JSONL file at `~/.openclaw/workspace/memory/episodic/YYYY-MM-DD.jsonl`. The entry schema includes: `id`, `timestamp`, `session_id`, `project`, `content`,

`tokens`, `promoted` (Boolean), and `cognitive_weight` $\in [-1, 1]$ (initialised to 0). **Cognitive Weight (CW)** is a per-entry scalar that accumulates evidence of retrieval quality over time: positive values indicate that the entry has contributed to successful tool executions; negative values indicate association with failures. CW starts at 0 (neutral) and is updated by the attribution loop (§3.2) after each agent session. It serves as the system’s long-term memory of *which memories have proven useful*, allowing the retrieval engine to prefer high-quality entries without human labelling. Entries are append-only (no in-place mutation) and project-scoped, preventing cross-project contamination. To support multi-agent orchestration (e.g., OpenClaw’s Lobster engine), episodic logs are **agent-private by default**: sub-agents write only to their own episodic ledger, preventing immediate context bleed. An orchestrator agent may be granted global read access via project-level configuration. System entries (prefixed `[system]`) are written but excluded from retrieval.

3.2 Phase 1b: Weighted Retrieval Engine

The scoring function follows Eq. (1):

$$S(q, m_i) = \mathbf{w}^\top \phi(q, m_i) \quad (1)$$

where $\phi = [\phi_{\text{sem}}, \phi_{\text{bm25}}, \phi_{\text{decay}}, \phi_{\text{cw}}, \phi_{\text{tier}}]$ and default weights $\mathbf{w}_0 = [0, 0.35, 0.25, 0.25, 0.15]$.

BM25 signal. Okapi BM25 with $k_1=1.5$, $b=0.75$, normalised to $[0, 1]$. High BM25 scores (> 2.0) suppress time decay (decay bypass rule).

Time decay signal. $\phi_{\text{decay}} = e^{-\lambda \Delta t}$ with $\lambda = 0.05$ (half-life ≈ 14 days). Entries from semantically-relevant sessions bypass decay regardless of age.

Cognitive weight signal. $\phi_{\text{cw}} = (\text{CW} + 1)/2$, mapping $[-1, 1]$ to $[0, 1]$. Entries used in successful tool calls accumulate positive CW; entries linked to failures accumulate negative CW.

Tier boost signal. The tier multiplier $\mu_k \in \{1.0, 1.2, 1.4\}$ for episodic, semantic, and procedural tiers is applied as an *additive bonus* outside the dot product: $S = \mathbf{w}^\top \phi + w_{\text{tier}}(\mu_k - 1)$, where ϕ contains only the first four signals. This preserves the interpretation of $\mathbf{w}^\top \phi$ as a relevance score and treats the tier bonus as a separable promotion incentive.

Hyperparameter justification. The default weights $\mathbf{w}_0 = [0, 0.35, 0.25, 0.25, 0.15]$ are initialised to reflect our *a priori* signal reliability ordering: BM25 is the most direct measure of query relevance and receives the highest weight (0.35); time decay and cognitive weight are both informative proxies but noisier, so they share a lower weight (0.25 each); tier boost is a structural incentive rather than a relevance signal and receives the smallest weight (0.15). These are deliberately conservative starting values—the PPO trainer is designed to learn away from them.

$k_1=1.5$, $b=0.75$ are the canonical Okapi BM25 defaults of Robertson et al. (1994), widely used as strong baselines across IR tasks (Lin, 2021). The decay parameter $\lambda=0.05$ (half-life ≈ 14 days) is motivated by the observation that knowledge-work context windows typically span one to two weekly sprint cycles; entries older than two weeks are unlikely to be directly relevant to the current task. The BM25 bypass threshold of 2.0 is a conservative gate: a raw BM25 score above 2.0 requires at least two rare shared terms ($\text{IDF} > 1$) in a short document, indicating a strong lexical match that should override recency preference. The tier multipliers $\mu_k \in \{1.0, 1.2, 1.4\}$ are deliberately modest (20% and 40% boosts) to avoid tier dominance overriding relevance scores; they act as a tiebreaker, not a reranker. All hyperparameters are candidates for learning via the PPO trainer and may diverge from these defaults as the episode pool diversifies.

Two-stage retrieval. Stage 1 indexes the semantic tier with BM25 and extracts the top-5 relevant session IDs. Stage 2 loads episodic entries scoped only to those sessions and scores them with the full formula. This reduces the retrieval pool from all stored entries to a focused subset, preventing context overflow.

Token efficiency. Semantic pre-population reduces the number of facts indexed per question from ~ 509 (heuristic) to ~ 3.1 (LLM-extracted), a $164\times$ reduction. The smaller, higher-precision semantic index has two practical benefits beyond accuracy: (a) Stage 1 BM25 lookup is faster and more discriminating, and (b) the facts injected into the generator’s context consume fewer tokens while carrying more signal. The pre-population cost is low: ~ 500 API calls and approximately \$0.05 total for 500 questions at DeepSeek-V4-Flash pricing, as a one-time offline preprocess-

ing step. For live agent deployments, the consolidation daemon (§3.3) performs equivalent LLM-quality extraction continuously during operation, so the offline preprocessing step is only required for batch benchmark evaluation.

3.3 Ablation Study

We design a systematic ablation to isolate the contribution of each MEMTIER component across three axes: (a) **component removal**, removing one signal or pipeline stage; (b) **retrieval budget** sensitivity (k , the number of injected entries); and (c) **token injection budget** sensitivity. Table 2 presents all results ($N=500$, all runs complete).

Component importance. The ablation reveals a clear hierarchy. **Semantic pre-population** is the dominant contributor: removing it costs -0.128 Acc and reduces F1 by a $51\times$ factor. **Two-stage scoping** is the second most important component (-0.038 Acc): without session pre-selection the episodic BM25 pool becomes noisy across all 53 sessions. The individual signals—time decay, cognitive weight, tier boost—each contribute approximately equally ($\Delta\text{Acc } -0.014$ to -0.016) and their contributions are additive and non-redundant.

Optimal retrieval budget: $k = 2$. $k = 2$ outperforms the default $k = 4$ (Acc=0.402 vs. 0.380, $+0.022$). With high-precision semantic pre-population active, the top-2 entries are already highly relevant; additional entries introduce noise that a 7B generator cannot reliably filter. $k = 8$ partially recovers (0.394), suggesting a non-monotonic relationship between k and accuracy for small generators. For edge deployments with constrained generators, $k = 2$ is recommended.

Token budget is a real constraint. Increasing from 300 to 600 tokens yields Acc=0.412, F1=0.427 ($+0.032$), confirming that some answers require more injected context. Reducing to 150 tokens costs only -0.006 , demonstrating graceful degradation—useful when inference budget is severely constrained. We recommend 600 tokens as the default when context budget permits.

3.4 LoCoMo

We also evaluate on LoCoMo (Maharana et al., 2024), a 10-session conversational benchmark with 200 QA pairs. As Table 3 shows, MEMTIER

Table 1: LongMemEval-S accuracy by question type (N=500). MEMTIER BM25 only uses raw episodic retrieval; MEMTIER + semantic adds DeepSeek-V4-Flash fact pre-population. Taxonomy follows Wu et al. (2025).

Type	N	Full-ctx-7B	MEMTIER BM25 only	MEMTIER + semantic
Single-session user	70	0.057	0.543	0.686
Single-session assistant	56	0.107	0.464	0.732
Knowledge update	78	0.000	0.346	0.436
Temporal reasoning	133	0.105	0.203	0.316
Multi-session	133	0.008	0.060	0.180
Single-session preference	30	0.000	0.000	0.067
Overall Acc	500	0.050	0.252	0.382
Overall F1	500	0.054	0.142	0.412

Table 2: Ablation study on LongMemEval-S (N=500). Full system: Acc=0.382, F1=0.412 ($k = 4$, 300 tokens). Each row removes or modifies one component. Δ Acc relative to full system.

Configuration	Acc	F1	Δ Acc
<i>Reference points</i>			
Full MEMTIER + semantic (ours)	0.382	0.412	—
MEMTIER + semantic + PPO-learned weights	0.382	0.412	± 0.000
MEMTIER BM25 only (— semantic tier)	0.252	0.142	−0.128
Full-context-7B (no retrieval)	0.050	0.054	−0.330
<i>Signal removal (weight = 0, renormalised)</i>			
— time decay ($w_{\text{decay}} = 0$)	0.394	0.409	−0.014
— cognitive weight ($w_{\text{cw}} = 0$)	0.396	0.409	−0.016
— tier boost ($w_{\text{tier}} = 0$)	0.394	0.409	−0.014
— two-stage scoping	0.342	0.380	−0.038
<i>Retrieval entries k</i>			
$k = 1$	0.386	0.397	−0.006
$k = 2$ (optimal)	0.402	0.414	+0.022
$k = 4$ (default)	0.382	0.412	—
$k = 8$	0.394	0.410	+0.014
<i>Token injection budget</i>			
150 tokens	0.374	0.397	−0.006
300 tokens (default)	0.382	0.412	—
600 tokens (recommended)	0.412	0.427	+0.032

Table 3: LoCoMo results. Both systems use the same context-stuffed prompt; the memory architecture is irrelevant.

System	F1	R@1
SimpleMem [†]	0.432	—
MemGPT [†]	0.218	—
OpenClaw-Default (ours)	0.125	0.105
MEMTIER Phase 1b (ours)	0.120	0.100

and the unmodified baseline score identically (Δ F1 = −0.005).

This null result is informative: LoCoMo inserts the full conversation into context at query time, making the memory architecture irrelevant. We concur with the recommendation in Wu et al. (2025) that LongMemEval is the appropriate benchmark for evaluating memory *storage and retrieval*, while LoCoMo tests in-context comprehension. We report LoCoMo for completeness and

Table 4: Retrieval weight vector before and after PPO training (15 seed episodes). Seed data has near-uniform CW, producing near-zero advantage estimates and near-zero gradients by design.

Signal	w_0	w_{PPO}
Semantic (w_{sem})	0.000	0.000
BM25 (w_{bm25})	0.350	0.350
Time decay (w_{decay})	0.250	0.250
Cognitive weight (w_{cw})	0.250	0.250
Tier boost (w_{tier})	0.150	0.150
Mean reward (r)	—	0.303

comparison with prior work.

3.5 PPO Weight Trainer: Infrastructure Validation

Table 4 shows the weight vector after 15 training episodes on seed data.

The weights do not diverge in this initial 15-

episode run. Post-review analysis identifies two mathematical causes: **(1) Circular reward trap:** the CW-based reward is derived from Jaccard attribution, which is itself a proxy—the PPO is optimising a proxy of a proxy with no ground-truth signal. **(2) Zero-variance trap:** with 15 seed episodes of near-uniform CW, the advantage $A_t = r_t - \bar{r} \approx 0$ for every episode, so the gradient update is identically zero.

The fix, implemented and described here for camera-ready completion, replaces CW-based reward with direct task success:

$$r_t = \begin{cases} +1.0 & \text{if } \text{em}(\hat{a}_t, a^*) = 1 \\ -1.0 & \text{otherwise} \end{cases} \quad (2)$$

where a^* is the LongMemEval-S gold answer. This constitutes true credit assignment: the weight vector $\mathbf{w}$ that produced a correct retrieval receives a direct positive signal. Exploration is forced by initialising $\sigma = 0.15$ (vs. near-zero previously), ensuring the advantage has real variance and gradients flow. Training runs on 100 stratified LongMemEval-S questions (4 epochs, batch=16) and is expected to produce meaningful weight divergence; results will be included in the camera-ready submission.

4 Discussion

Three-layer invariance: architecture is the ceiling. Our evaluation reveals invariance across three axes, converging on a single conclusion. **Generator invariance:** DeepSeek-V4-Flash (284B MoE, 13B active) yields Acc=0.234 vs. Qwen2.5-7B Acc=0.252—statistically identical with the same per-category profile. **Weight invariance:** PPO-learned weights after task-success training (100 questions, 4 epochs, +1/−1 reward) yield Acc=0.382, F1=0.412—identical to default weights—because ± 0.03 reweighting cannot alter BM25-dominated retrieval rankings. **Together:** neither the generator, nor the retrieval weights, nor their interaction determines performance—the BM25 retrieval *architecture* is the binding constraint. Better models and better weights both fail when the retrieval stage cannot surface multi-session evidence (0.180) or resolve temporal references (0.316). This directly motivates recall-first retrieval and dense/hybrid scoring as the necessary Phase 3 components.

The benchmark selection problem. Our Lo-CoMo null result replicates an often-overlooked

issue: benchmarks that make conversations available in context at query time cannot distinguish memory architectures. We recommend that future work on agent memory systems report LongMemEval-S as the primary metric, with Lo-CoMo as a secondary in-context-comprehension baseline.

Limitations. (1) **Attribution path:** SGLang logprob attribution is code-complete but blocked by local hardware constraints on the evaluation machine; the lexical Jaccard fallback used in production is a coarse proxy. (2) **RL weight dominance:** While direct task-success reward (+1/−1) enabled PPO gradients to flow and weights to shift, BM25’s unbounded scoring heavily dominates the bounded signals (decay, CW) in the linear combination, masking the RL’s impact on final ranking. Future iterations must strictly normalise BM25 or transition to dense retrieval before learned weights can provide marginal gains. (3) **Relation extraction:** Heuristic KV-pattern extraction produces coarse labels (*e.g.*, `mentioned_in`); a fine-grained NLP extractor would further improve semantic tier quality.

5 Conclusion

We presented MEMTIER, a five-phase memory architecture for long-running autonomous agents that addresses context collapse, compaction discontinuity, structural blindness, and the absence of an attribution loop. Evaluated on the full 500-question LongMemEval-S benchmark, MEMTIER with semantic pre-population achieves **Acc=0.382** and **F1=0.412** with a 7B generator on edge hardware—a +33 percentage point improvement over the full-context baseline. On single-session recall, it achieves 0.686–0.732, exceeding the paper’s RAG BM25 GPT-4o baseline (0.560).

The central diagnostic finding is a **three-layer invariance:** neither scaling the generator to a 284B MoE model, nor dynamically tuning retrieval weights via PPO with true credit assignment, significantly altered the performance ceiling dictated by the BM25 retrieval architecture. This empirically demonstrates that high-precision, structurally isolated memory (episodic vs. semantic) is the primary driver of agentic long-horizon success, while simultaneously demonstrating that legacy linear-combination retrieval is the current bottleneck. Future work will transition to recall-first dense retrieval to solve multi-session syn-

thesis (0.180), implement absolute date resolution for temporal reasoning (0.316), and activate the SGLang logprob attribution path for higher-fidelity RL credit assignment.

Limitations

(1) **Attribution path:** SGLang logprob attribution is code-complete but blocked by local hardware constraints on the evaluation machine; the lexical Jaccard fallback used in production is a coarse proxy. (2) **RL weight dominance:** While direct task-success reward (+1/ − 1) enabled PPO gradients to flow and weights to shift, BM25’s unbounded scoring heavily dominates the bounded signals (decay, CW) in the linear combination, masking the RL’s impact on final ranking. Future iterations must strictly normalise BM25 or transition to dense retrieval before learned weights can provide marginal gains. (3) **Relation extraction:** Heuristic KV-pattern extraction produces coarse labels (*e.g.*, `mentioned_in`); a fine-grained NLP extractor would further improve semantic tier quality.

6 Conclusion

We presented MEMTIER, a five-phase memory architecture for long-running autonomous agents that addresses context collapse, compaction discontinuity, structural blindness, and the absence of an attribution loop. Evaluated on the full 500-question LongMemEval-S benchmark, MEMTIER with semantic pre-population achieves **Acc=0.382** and **F1=0.412** with a 7B generator on edge hardware—a +33 percentage point improvement over the full-context baseline. On single-session recall, it achieves 0.686–0.732, exceeding the paper’s RAG BM25 GPT-4o baseline (0.560).

The central diagnostic finding is a **three-layer invariance**: neither scaling the generator to a 284B MoE model, nor dynamically tuning retrieval weights via PPO with true credit assignment, significantly altered the performance ceiling dictated by the BM25 retrieval architecture. This empirically demonstrates that high-precision, structurally isolated memory (episodic vs. semantic) is the primary driver of agentic long-horizon success, while simultaneously demonstrating that legacy linear-combination retrieval is the current bottleneck. Future work will transition to recall-first dense retrieval to solve multi-session synthesis (0.180), implement absolute date resolution

for temporal reasoning (0.316), and activate the SGLang logprob attribution path for higher-fidelity RL credit assignment.

Acknowledgments

All experiments were conducted on a consumer laptop with a 6GB GPU. We thank the OpenClaw open-source community for maintaining the runtime used in all experiments. The authors used Claude (Anthropic) for writing assistance and code generation support; all scientific claims and results are the authors’ own.

References

- Anonymous. 2025a. A-MEM: Adaptive memory for LLM agents. *arXiv preprint*. See paper for full reference.
- Anonymous. 2025b. Memory-R1: Reinforcement learning for conversational memory. *arXiv preprint*. See paper for full reference.
- Anonymous. 2025c. MIRIX: Multi-instance retrieval index for LLM agents. *arXiv preprint*. See paper for full reference.
- Anonymous. 2026. AgentWarden: RL-based adaptive capability governance for AI coding agents. *NeurIPS 2026 Agent Safety Workshop (under review)*.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. [Retrieval-augmented generation for knowledge-intensive NLP tasks](#). In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Jimmy Lin. 2021. [A few brief notes on DeepImpact, COIL, and a conceptual framework for information retrieval techniques](#). *arXiv preprint arXiv:2106.14807*.
- Hao Liu and 1 others. 2026. SimpleMem: Efficient lifelong memory for LLM agents. *arXiv preprint*.
- Adyasha Maharana, Divyansh Das, Sergey Tulyakov, Mohit Bansal, Franck Dernoncourt, and Yuwei Fang. 2024. [Evaluating very long-term conversational memory of LLM agents](#). *arXiv preprint arXiv:2402.17753*.
- OpenClaw Contributors. 2026. [OpenClaw: Personal AI assistant framework](#). Open-source agent runtime. Issues #33406, #62488 referenced in this paper.
- Charles Packer, Vivian Fang, Shishir G. Patil, Kevin Wooders, and Joseph E. Gonzalez. 2023. [MemGPT: Towards LLMs as operating systems](#). *arXiv preprint arXiv:2310.08560*.

- Stephen Robertson, Steve Walker, Susan Jones, Micheline Hancock-Beaulieu, and Mike Gatford. 1994. [Okapi at TREC-3](#). In *Proceedings of the Third Text REtrieval Conference (TREC-3)*, pages 109–126.
- Wei Sun and 1 others. 2026. H-MEM: Hierarchical memory for high-efficiency long-term reasoning in LLM agents. *arXiv preprint*.
- Di Wu and 1 others. 2025. [LongMemEval: Benchmarking chat assistants on long-term interactive memory](#). *arXiv preprint arXiv:2501.08956*.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. [ReAct: Synergizing reasoning and acting in language models](#). In *International Conference on Learning Representations (ICLR)*.